

Recommendation

Use Presentation Orchestrator when AI slide generation must become decision preparation

The skill is not a slide beautifier. It is a consulting-style harness for decks where a room must approve, choose, fund, stop, or assign.

Decision ask

Pilot it on high-stakes business decks that need a clear POV, source-aware evidence, rejected paths, red-team rehearsal, and a specific next action.

Use for	Avoid for	Core proof	Human role	Output
approval, strategy, GTM, roadmap	quick cosmetic formatting	case-led deck packet	judge assumptions and risk	PPTX/PDF + speaker notes

Main claim: the skill improves the presenter’s decision readiness, not merely slide aesthetics.

The skill is worth using only when it can prove a stronger decision deck

Recommendation


Use the skill as a deck-strategy partner for high-stakes decisions; do not position it as a generic PPT generator.

Proof point	What it changes
1. Decision spine	Turns rough input into an issue question, route card, and explicit ask.
2. Evidence discipline	Separates user facts, source-backed facts, assumptions, inference, and gaps.
3. Rehearsal loop	Forces sponsor, skeptic, operator, and evidence-review objections before final output.
4. Verified artifact path	This deck is generated reproducibly and validated as PPTX/PDF; package checks pass.

Consulting-team critique absorbed

The prior version explained the process. This version uses one synthetic GTM case to prove judgment: selected method, rejected paths, evidence gaps, red-team changes, and readiness gates.


A messy GTM request becomes a decision question the room can answer



Audience	Decision	Time box	Evidence state	Output
growth leadership	approve / defer	6 weeks	partial and mixed	decision deck

Source: Illustrative synthetic case; no market facts claimed

The skill chooses the decision job before it chooses the framework




The evidence ledger turns weak signals into explicit confidence and next checks

Claim	Source type	Conf.	Gap / risk	Slide implication
Vertical segment shows interest	user notes	Medium	sample biased toward warm prospects	ask for pilot, not rollout
Current team can execute pilot	assumption	Low	owner capacity not confirmed	name owner before approval
Channel economics may work	inference	Low	CAC/payback unproven	define 2 evidence gates
Full rollout is premature	inference	High	proof is not strong enough	reject scale-now option
Defer loses learning	assumption	Medium	opportunity cost unknown	time-box the pilot

What would change our mind


If pilot CAC exceeds threshold, owner capacity is unavailable, or pipeline quality fails, the recommendation flips from pilot to defer.

The selected framework changes the decision: pilot beats scale-now and defer



Source: Illustrative synthetic case; no external market claim

The deck packet turns analysis into five action-title slides



Presenter packet

The output is not only slide titles. Each page carries a proof standard, speaker note, appendix cue, and known objection so the presenter can rehearse the decision path.

Quality bar: if a slide cannot change the decision, support the ask, or expose a risk, it moves to appendix or disappears.

The rehearsal loop improves the recommendation before the room reacts

Persona	Tough question	Revision made
CFO	What evidence justifies funding?	Funding ask changed from rollout to time-boxed pilot with kill criteria.
Operator	Who owns Monday execution?	Added owner confirmation as approval condition.
Sales leader	Will this distract from current pipeline?	Added capacity check and pilot scope limit.
Evidence reviewer	Which claims are assumptions?	Ledger now labels source type and confidence on every major claim.

Why this matters

The red-team loop is valuable only if it changes the output. Here it changed the ask, the approval conditions, and the confidence language.

The orchestrated deck makes decision-readiness gates explicit

	Generic prompt risk		Orchestrated output
Decision ask	unclear or buried	—	approve gated pilot / defer / reject rollout
Framework use	decorative model	—	issue tree selected; BMC and Porter rejected
Evidence discipline	facts and assumptions mixed	—	source type, confidence, gaps, appendix cue
Alternatives	single-path story	—	scale now, gated pilot, defer compared
Objections	left for Q&A	—	CFO/operator/sales/evidence objections handled
Action	next steps vague	—	owner, proof gates, kill criteria, next review
Proof standard This is a designed readiness comparison, not a claim that every generic AI output fails. A real pilot should test before/after quality on live decks.			

Use the skill for decision decks; keep humans accountable for facts and judgment

Use it for


strategy, GTM, product, board, investment, roadmap, workshop readout

Do not use it for

pure formatting, unverified claims, confidential upload without approval, legal or financial certainty

Human review required


facts, source permission, assumptions, recommendations, risk tolerance, final business decision



Operating principle

The skill can structure the work and expose gaps; it cannot manufacture proof, replace source review, or own the business decision.

The artifact path is reproducible and current package checks pass



What this proves

The deck artifact, source references, and public skill package are reproducible and locally validated.

It does not prove every future recommendation is true; facts, source permissions, and business judgment still require human review.

Scope of proof

Technical validation proves package integrity and renderability. Advisory quality is validated only through evidence review and live before/after use.

Decision readiness is the success measure, not slide count

Final gate

Use the skill when the cost of a weak deck is a weak decision.

Unclear ask, late objections, hidden assumptions, or no accountable next owner.

1. Decision ask

approve / reject / test is explicit

2. Evidence standard

source type, confidence, and gaps are visible

3. Human accountability

owner, proof gates, and final judgment remain named

Pilot rubric: compare 3 real decks to baseline; require $\geq 4/5$ on ask, evidence, alternatives, owner, and rehearsal changes.

Ship as a public demo artifact; use live before/after evidence before claiming business-outcome improvement.